

# Task Description

The goal of this exercise is to model the relationship between weather observations and the amounts of cyclists in the city of Vienna.

To investigate a potential relationship, we will use two datasets:

- tri-daily weather reports from 2009 - 2023 of Vienna.
- daily cyclists in Vienna from 2021 - 2022.

In this exercise, you will

- use `pandas` to read, prepare and transform data,
- use `matplotlib` to visually analyse data,
- use `scikit-learn` to build prediction models.

The data to be used can be found in the provided paths

To complete this exercise, you will have to:

- prepare the data, which (at minimum) involves the following:
  - load and prepare the data
  - handling missing values
  - handling outliers
  - temporal alignment of the two datasets
- analyse the data:
  - compare descriptive statistics
  - visually investigate the raw data to gain an understanding of the data identify patterns, outliers etc.,

**As this is a practice notebook, we are showing you all hidden tests, during the exams they will not be visible**

**Keep in mind that the exam also tests different aspects of Data Processing then this notebook does, and contains more programming and multiple choice questions**

```
# DO NOT MODIFY OR COPY THIS CELL!!
# Note: The only imports allowed are Python's standard library,
# pandas, numpy, scipy, matplotlib, seaborn and scikit-learn
import numpy as np
import pandas as pd
import glob
import os
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns
import datetime
import pickle
import typing
```

```

data_path = os.path.join(os.environ["HOME"], "shared", "194.192-2025W", "data", "cyclists")
weather_data_path = os.path.join(data_path, 'weather')
cyclists_data_path = os.path.join(data_path, 'cyclists')

```

## Data Loading

As a first step, implement the method `load_weather_data()`, which should read all individual (yearly) datasets from the csv files in `data/weather/` into a single `pd.DataFrame` and return it.

- make sure that you load all the data (2009-2023, 15 years)
- split the tri-daily and daily data (tri-daily data has \_7h, \_14h, and \_19h suffixes for column headers), and convert the tri-daily data from a wide to a long format (use pandas' `wide_to_long` or `melt` functions). Introduce a new `hours` column which's values should be taken from the column suffixes.
- make sure all columns are appropriately typed (numeric values -> float, countables, i.e. days -> int, etc.)! Especially the `date` column! See `datetime` and `pandas.Timestamp`!
- from the `date` column create `year`, `month`, `day` columns. Use Pandas built-in `datetime` handling features.
- from the wide to long transform, you should have an `hour` column with the 7, 14, or 19 hours values.
- create a `MultiIndex` from the date columns with the following hierarchy: `year - month - day - hour` (make sure to label them accordingly)

### Hints:

- LOOK at the data in the original files
- It is advisable not to append each data set individually, but to read each data frame, store it into a list and combine them once at the end.
- Note that for the `precip` data column you will get an unexpected (object) datatype. For this task it is ok to leave it like that, it is due to special values, see next chapters!
- You will find similar names, in the disruption column, you are free to combine them to reduce the number of different disruptions.
- Your resulting data frame should look as follows, with temperature in Celsius, air pressure in hecto Pascal, skyCover on a scale from 1-10, humidity in percent, windDir in compass directions, windBeauf in Beaufort and precip in millimeters.: alt text

```

def load_weather_data(weather_data_path: str) ->
    typing.Tuple[pd.DataFrame, pd.DataFrame]:
    """
        Load all weather data files and combine them into a single Pandas DataFrame.
        Split the tri-daily data from the daily data.
        For the tri-daily data create a new hour column using the indicated hour in the column names.
        Add a hierarchical index (year, month, day, hour).
        For the daily-only data also add a hierarchical index (year,

```

```
month, day).
```

*Parameters*

```
-----  
weather_data_path: path to directory containing weather data CSV  
files
```

*Returns*

```
-----  
weather_data: data frame containing the tri-daily (hours) weather  
data  
weather_data_daily: data frame containing the daily weather data  
(e.g. precip, precipType, etc.)  
"""
```

*# YOUR CODE HERE*

```
csv_files = glob.glob(os.path.join(weather_data_path, "*.csv"))  
df_list = [pd.read_csv(file, sep=";") for file in csv_files]  
combined_df = pd.concat(df_list, ignore_index=True)  
  
combined_df["date"] =  
pd.to_datetime(combined_df["date"], format="%d.%m.%Y")  
combined_df["year"] = combined_df["date"].dt.year  
combined_df["month"] = combined_df["date"].dt.month  
combined_df["day"] = combined_df["date"].dt.day  
  
weather_data = combined_df  
stubs =  
["airPressure", "skyCover", "temp", "hum", "windDir", "windBeauf"]  
weather_data = pd.wide_to_long(weather_data, stubs,  
i="date", j="hour", sep="_", suffix=r"\d+h")  
weather_data = weather_data.reset_index()
```

*#type conversion*

*# make hour an integer*

```
weather_data["hour"] = (  
    weather_data["hour"]  
    .astype(str)  
    .str.extract(r"(\d+)", expand=False)  
    .astype(int))
```

*# make sure year, month, day are ints too*

```
weather_data["year"] = weather_data["year"].astype(int)  
weather_data["month"] = weather_data["month"].astype(int)  
weather_data["day"] = weather_data["day"].astype(int)
```

```
weather_data = weather_data.set_index(["year", "month", "day"],
```

```

"hour"]).sort_index()

    weather_data_daily = weather_data.drop(labels=stubs, axis=1)
    weather_data =
weather_data.drop(labels=["precip", "precipType"], axis=1)

    col = weather_data_daily.pop("precip")
    weather_data_daily["precip"] = col

#fahrenheit conversion
weather_data["temp"] = (weather_data["temp"] - 32) * 5/9

    return weather_data, weather_data_daily

# DO NOT MODIFY OR COPY THIS CELL!
weather_data, daily_weather_data =
load_weather_data(weather_data_path)
# print first couple of rows:
print('hourly weather data:')
display(weather_data.head())
print('\ndaily weather data:')
display(daily_weather_data.head())

hourly weather data:



| windDir \ |       |     |            |            |        |  |    |            |    |
|-----------|-------|-----|------------|------------|--------|--|----|------------|----|
| year      | month | day | hour       |            |        |  |    |            |    |
| 2009      | 1     | 7   | 2009-01-01 |            | 999.7  |  | 10 | -20.277778 | 79 |
| W         |       | 14  | 2009-01-01 |            | 998.8  |  | 5  | -17.833333 | 71 |
| NW        |       | 19  | 2009-01-01 |            | 1000.7 |  | 10 | -18.777778 | 72 |
| NW        |       | 2   | 7          | 2009-01-02 | 999.6  |  | 10 | -19.444444 | 67 |
| NaN       |       | 14  | 2009-01-02 |            | 998.5  |  | 9  | -18.500000 | 66 |
| W         |       |     |            |            |        |  |    |            |    |
|           |       |     |            | windBeauf  |        |  |    |            |    |
| year      | month | day | hour       |            |        |  |    |            |    |
| 2009      | 1     | 7   |            | 2          |        |  |    |            |    |
|           |       | 14  |            | 2          |        |  |    |            |    |
|           |       | 19  |            | 2          |        |  |    |            |    |
|           | 2     | 7   |            | 0          |        |  |    |            |    |
|           |       | 14  |            | 2          |        |  |    |            |    |



daily weather data:

```

year	month	day	hour	date	precipType	precip
2009	1	7	14	2009-01-01	Nan	0
			19	2009-01-01	Nan	0
	2	7	2009-01-02	snow	traces	
		14	2009-01-02	snow	traces	

```

# DO NOT MODIFY OR COPY THIS CELL!!
# TESTS: dimensions should be like this:
assert weather_data.shape[0] == 16434
assert weather_data.shape[1] >= 7

#### TESTS
#check for Fahrenheit Conversion
assert weather_data[weather_data['date'] >= '2023-01-01']['temp'].min() < 0

#### TESTS
assert pd.date_range(start = '2009-01-01', end = '2023-12-31').difference(weather_data.date).empty
assert weather_data.set_index('date').index.difference(pd.date_range(start = '2009-01-01', end = '2023-12-31', freq='1D')).empty

#### TESTS
# check if all dates are present
index_dtypes = weather_data.index.dtypes

assert pd.api.types.is_integer_dtype(index_dtypes.iloc[0])
assert pd.api.types.is_integer_dtype(index_dtypes.iloc[1])
assert pd.api.types.is_integer_dtype(index_dtypes.iloc[2])
assert pd.api.types.is_integer_dtype(index_dtypes.iloc[3])

assert pd.api.types.is_datetime64_any_dtype(weather_data['date'].dtype)
assert pd.api.types.is_float_dtype(weather_data['temp'].dtype)
assert pd.api.types.is_float_dtype(weather_data['airPressure'].dtype)
assert pd.api.types.is_integer_dtype(weather_data['hum'].dtype)
assert pd.api.types.is_integer_dtype(weather_data['skyCover'].dtype)
assert pd.api.types.is_integer_dtype(weather_data['windBeauf'].dtype)
assert pd.api.types.is_string_dtype(weather_data['windDir'].dtype)

assert pd.api.types.is_string_dtype(daily_weather_data['precip'].dtype)
assert pd.api.types.is_string_dtype(daily_weather_data['precipType'].dtype)

```

In which month was the average temperature the lowest?

- Implement the function below to find the answer!
- Find the respective entry/entries using pandas!!

```
def get_lowest_average_temp(data_frame:pd.DataFrame):
    year = 0
    month = 0

    monthly = (
        data_frame.dropna(subset=["temp"])
        .groupby(["year", "month"])["temp"]
        .mean().reset_index()
    )
    i = monthly["temp"].idxmin()
    year = int(monthly.loc[i, "year"])
    month = int(monthly.loc[i, "month"])

    return year, month

# DO NOT MODIFY OR COPY THIS CELL!!
low_num_year, low_num_month = get_lowest_average_temp(weather_data)
print(f"Month {low_num_month}, of year {low_num_year} has the lowest
average temperature!")

Month 1, of year 2017 has the lowest average temperature!

# hidden tests, DO NOT MODIFY OR COPY THIS CELL!!
### Tests
assert low_num_year == 2017
assert low_num_month == 1
```

## Task: Implement the `load_cycling_data` Method

In this exercise, you will write a function to **load and prepare cycling traffic data from Vienna**. The goal is to read the dataset, enrich it with useful time-related information, and return a well-structured DataFrame.

Steps you should implement:

1. **Load the dataset**
  - Read the CSV file `cyclists.csv` from the provided `cyclist_data_path`.
2. **Add weekday information**
  - Create a new column `day_of_week` that contains the weekday name (e.g., *Monday, Tuesday*).
3. **Reorganize the DataFrame index**
  - Sort the data by `year`, `month`, and `day`.
  - Set these columns as a **multi-index** (`year`, `month`, `day`) for easier time-based access.
4. **Return the processed data**

- The function should return the cleaned and enriched DataFrame with cycling traffic data.

```
def load_cycling_data() -> pd.DataFrame:

    # YOUR CODE HERE
    file=cyclists_data_path+"/cyclists.csv"
    df = pd.read_csv(file,sep=",")
    df["date"] = pd.to_datetime(df["date"])
    df["day_of_week"] = df["date"].dt.day_name()
    df["date"] = pd.to_datetime(df["date"],format="%d.%m.%Y")
    df["year"] = df["date"].dt.year
    df["month"] = df["date"].dt.month
    df["day"] = df["date"].dt.day
    df = df.set_index(["year", "month", "day"]).sort_index()
    data = df

    return data

data_cyclists = load_cycling_data()
display(data_cyclists)
```

year	month	day	date	number	day_of_week
2021	1	1	2021-01-01	412	Friday
		2	2021-01-02	648	Saturday
		3	2021-01-03	707	Sunday
		4	2021-01-04	1006	Monday
		5	2021-01-05	1198	Tuesday
...	...	...	...	...	...
2022	12	27	2022-12-27	1091	Tuesday
		28	2022-12-28	1245	Wednesday
		29	2022-12-29	1102	Thursday
		30	2022-12-30	1159	Friday
		31	2022-12-31	782	Saturday

[730 rows x 3 columns]

This method calculates the **average number of cyclists** for a given season and year from the dataset.

- It first ensures the **date** column is in datetime format.
- Seasons are defined by months:
  - **Spring:** March–May
  - **Summer:** June–August
  - **Autumn:** September–November

- **Winter:** December (previous year) + January–February (current year)
- The method filters the DataFrame for the specified `season` and `year`.
- It then computes the mean of the `number` column (cyclist counts) and returns the rounded result.

```
def average_cyclists_per_season(df: pd.DataFrame, season: str, year: int) -> int:
    # YOUR CODE HERE
    df = df.copy()
    df["date"] = pd.to_datetime(df["date"])
    df["month"] = df["date"].dt.month
    df["year"] = df["date"].dt.year

    seasons = {
        "spring": [3, 4, 5],
        "summer": [6, 7, 8],
        "autumn": [9, 10, 11],
        "winter": [12, 1, 2],
    }

    months = seasons[season]

    if seasons == "winter":
        data = df[((df["year"] == year - 1) & (df["month"] == 12)) | ((df["year"] == year) & (df["month"].isin([1, 2])))]
    else:
        data = df[(df["year"] == year) & (df["month"].isin(months))]

    return int(data["number"].mean())

## Test
assert average_cyclists_per_season(data_cyclists, "summer", 2022)==3908
```

## Data Cleaning

### Temperature outliers

First we want to take a closer look at the temperature values. Check if we can identify some obvious outliers and come up with a strategy to handle/fix them.

In order to do so you will have to:

- Plot the temperature curve over time and a histogram of temperature values to identify possible outliers
- Plot a zoomed in version of individual outliers to get a better understanding what's happening
- Devise a strategy to get rid of outliers

```

def plot_value_series(df:pd.DataFrame, column:str) -> None:
    """
    Plot the values in column in data frame df
    """
    df.plot(kind='line', y=column, figsize=(15,3))

def plot_temp_analysis(df: pd.DataFrame) -> None:
    """
    Create two plots:
    1) Temperature values over time for the whole dataframe
    2) A histogram for temperature values.
        Choose appropriate bins enabling you to identify outliers!
    Parameters
    -----
    df: data frame containint the temperature values (temp) with
    potential outlier
    """
    x = pd.to_datetime(df["date"])

    # 1) Plot temperature over time
    plt.figure(figsize=(15, 3))
    plt.plot(x, df["temp"], linewidth=1)
    plt.title("Temperature Over Time")
    plt.xlabel("Time")
    plt.ylabel("Temperature")
    plt.grid(True)
    plt.tight_layout()
    plt.show()

    # 2) Histogram for temperature values
    temp_values = df["temp"].dropna()
    bins = int(np.sqrt(len(temp_values))) # simple bin rule: sqrt of
    sample size

    plt.figure(figsize=(10, 4))
    plt.hist(temp_values, bins=bins, color="skyblue",
    edgecolor="black")
    plt.title("Temperature Distribution (Histogram)")
    plt.xlabel("Temperature")
    plt.ylabel("Frequency")
    plt.grid(axis="y", alpha=0.75)
    plt.tight_layout()
    plt.show()

    # 3) Zoomed in
    start = pd.Timestamp("2021-01-01")
    end = pd.Timestamp("2022-03-15")
    ts = pd.to_datetime(x)


```

```

mask = (ts >= start) & (ts <= end)
plt.figure(figsize=(15, 3))
plt.plot(ts[mask], df.loc[mask, "temp"], linewidth=1)
plt.title("Temperature Zoom")
plt.xlabel("Time")
plt.ylabel("Temperature")
plt.grid(True)
plt.tight_layout()
plt.show()

plot_temp_analysis(weather_data)

```



### Remove temperature outliers

Implement the below function using the strategy you defined above to get rid of the temperature outliers

```

def handle_temp_outliers(noisy_data) -> pd.DataFrame:
    """
    Parameters
    -----
    noisy_data: data frame that contains temperature outliers ('temp' column)

    Returns
    -----
    cleaned_data: data frame with temperature outliers removed/handled
    """
    cleaned_data = noisy_data.copy()
    s = cleaned_data["temp"].astype(float)


    win = 31
    med = s.rolling(win, center=True, min_periods=win // 2).median()
    mad = (s - med).abs().rolling(win, center=True, min_periods=win // 2).median()
    sigma = 1.4826 * mad


    mask = (s - med).abs() > 3 * sigma
    cleaned_data["temp"] = s.where(~mask, med)

    return cleaned_data

# DO NOT MODIFY OR COPY THIS CELL!!
weather_data_cleaned = handle_temp_outliers(weather_data)
plot_temp_analysis(weather_data_cleaned)
print(weather_data_cleaned['temp'].describe())

```





```

count    16098.000000
mean     -9.312846
std      8.089690
min     -26.055556
25%     -14.722222
50%     -10.444444
75%     -6.166667
max      35.200000
Name: temp, dtype: float64

```

## Data inspection

Check the occurrence of the non-numeric values in the precipitation data. You can check the file `data/weather/weather_description_<year>.txt`, which might have additional clues what is going on.

Implement the function below and return a list of the non-numeric values that occur in the `precip` column of `daily_weather_data`. Make sure to only return every unique value once!

```

def get_non_numeric_precip_values(df:pd.DataFrame) -> typing.Set[str]:
    """
    Parameters
    -----
    df: data frame that contains non-numeric values in precip column

    Returns
    """

```

```

-----
non_numeric_values: list of unique non-numeric values.
Do not return duplicate values in the list!
"""
non_numeric_values = set()

# YOUR CODE HERE
df_filter = df["precip"].astype(str).str.replace('.', '',
1).str.isnumeric() == False
non_numeric_values = set(df.loc[df_filter,
"precip"].astype(str).unique())

return non_numeric_values

# DO NOT MODIFY OR COPY THIS CELL!!
non_numeric_values = get_non_numeric_precip_values(daily_weather_data)
print(f"\nnon-numeric values values: {non_numeric_values}")

non-numeric values values: {'traces'}

```

## Fix non-numeric values

Replace non-numeric values with some appropriate numerical values and convert the column to a more suitable data type. To get an idea, what appropriate values might be, check the file `data/weather/weather_description_<year>.txt` and the other numeric values in the `precip` column.

```

def fix_precip_values(df:pd.DataFrame) -> pd.DataFrame:
    """
    Parameters
    -----
    df: data frame that contains non-numeric values in precip column

    Returns
    -----
    ret_df: data frame with fixed precip values
    """
    ret_df = df.copy()

    # YOUR CODE HERE
    ret_df["precip"] =
    ret_df["precip"].astype(str).str.replace('traces', '0.05')
    ret_df["precip"] = ret_df["precip"].astype(float)

    return ret_df

daily_weather_data_fixed_precip =
fix_precip_values(daily_weather_data)

```

```

unique_values =
pd.unique(daily_weather_data_fixed_precip['precip'].values.ravel())
unique_values = np.sort(unique_values)
print(unique_values[:10])

[0.  0.05 0.1  0.2  0.3  0.4  0.5  0.6  0.7  0.8 ]

# test, DO NOT MODIFY OR COPY THIS CELL!!
assert
pd.api.types.is_float_dtype(daily_weather_data_fixed_precip['precip'].dtype), "precip should now be a float column!!"
assert daily_weather_data_fixed_precip.shape == daily_weather_data.shape, "do not remove or add rows!"

unique_values =
pd.unique(daily_weather_data_fixed_precip['precip'].values.ravel())
unique_values = np.sort(unique_values)
print(unique_values[:10])
assert np.isclose(unique_values[0], 0.0), "values for traces should be greater than 0.0mm!"
assert np.isclose(unique_values[2], 0.1), "values for traces should be smaller than 0.1mm!"

[0.  0.05 0.1  0.2  0.3  0.4  0.5  0.6  0.7  0.8 ]

# DO NOT MODIFY OR COPY THIS CELL!!
def fix_values_daily(data):
    """
    Parameters
    -----
    data: data frame containing missing values

    Returns
    -----
    complete_data: data frame not containing any missing values
    """
    complete_data = data.copy()
    complete_data = fix_precip_values(complete_data)

    return complete_data

def handle_outliers(data):
    """
    Parameters
    -----
    data: data frame containing outlier values

    Returns
    -----

```

```

complete_data: data frame not containing any outlier values
"""
complete_data = data.copy()
complete_data = handle_temp_outliers(complete_data)

return complete_data

# DO NOT MODIFY OR COPY THIS CELL!!
daily_weather_data_finished = fix_values_daily(daily_weather_data)
weather_data_finished = handle_outliers(weather_data)

weather_data_finished

```

				date	airPressure	skyCover	temp	hum	
windDir \	year	month	day	hour					
W	2009	1	1	7	2009-01-01	999.7	10	-20.277778	79
NW				14	2009-01-01	998.8	5	-17.833333	71
NW				19	2009-01-01	1000.7	10	-18.777778	72
NaN			2	7	2009-01-02	999.6	10	-19.444444	67
W				14	2009-01-02	998.5	9	-18.500000	66
...					...	...	...	...	...
...									
W	2023	12	30	14	2023-12-30	993.2	4	11.900000	58
W				19	2023-12-30	993.6	0	7.100000	78
SE			31	7	2023-12-31	988.8	10	3.200000	96
NaN				14	2023-12-31	985.1	10	3.700000	88
N				19	2023-12-31	984.1	9	2.900000	89
					windBeauf				
year	month	day	hour						
2009	1	1	7		2				
			14		2				
			19		2				
		2	7		0				
			14		2				
...					...				
...									
2023	12	30	14		3				
			19		2				

```
31    7          2  
14          0  
19          1
```

```
[16434 rows x 7 columns]
```

## Aggregate values

Aggregate the observations on a daily basis. Return a data frame with a hierarchical index (levels `year`, `month` and `day`) and the following daily aggregations as columns:

- `temp_dailyMin`: minimum of `temp`
- `temp_dailyMax`: max of `temp`
- `temp_dailyMean`: mean of `temp`
- `temp_dailyMedian`: median of `temp`
- `hum_dailyMin`: min of `hum`
- `hum_dailyMax`: max of `hum`
- `hum_dailyMean`: mean of `hum`
- `wind_dailyMean`: mean of `windBeauf`
- `wind_dailyMax`: max of `windBeauf`
- `wind_dailyMin`: min of `windBeauf`

Additionally merge the precipitation values from the `daily_weather_data` dataframe also into the newly created dataframe.

```
def aggregate_daily(hourly_data, daily_data):  
    # YOUR CODE HERE  
    hourly = hourly_data.reset_index().drop(columns=["hour"])  
    dd = daily_data.reset_index()  
  
    group_keys = ["year", "month", "day"]  
  
    daily = (  
        hourly.groupby(group_keys)  
        .agg(  
            temp_dailyMin=("temp", "min"),  
            temp_dailyMax=("temp", "max"),  
            temp_dailyMean=("temp", "mean"),  
            temp_dailyMedian=("temp", "median"),  
            hum_dailyMin=("hum", "min"),  
            hum_dailyMax=("hum", "max"),
```

```

        hum_dailyMean="hum", "mean"),
        wind_dailyMean="windBeauf", "mean"),
        wind_dailyMax="windBeauf", "max"),
        wind_dailyMin="windBeauf", "min"),
    )
    .reset_index()
)
daily["date"] = pd.to_datetime(daily[["year", "month", "day"]])
dd["date"] = pd.to_datetime(dd[["year", "month", "day"]])

merged = (
    daily.merge(dd[["date", "precip", "precipType"]], on="date",
how="outer")
    .drop(columns=["date"])
    .set_index(group_keys)
    .sort_index()
)
return merged
agg_daily_data = aggregate_daily(weather_data_finished,
daily_weather_data_finished)

agg_daily_data

```

			temp_dailyMin	temp_dailyMax	temp_dailyMean	\
year	month	day				
2009	1	1	-20.277778	-17.833333	-18.962963	
		1	-20.277778	-17.833333	-18.962963	
		1	-20.277778	-17.833333	-18.962963	
		2	-19.444444	-18.500000	-18.888889	
		2	-19.444444	-18.500000	-18.888889	
...			...	...	...	...
2023	12	30	7.100000	11.900000	9.933333	
		30	7.100000	11.900000	9.933333	
		31	2.900000	3.700000	3.266667	
		31	2.900000	3.700000	3.266667	
		31	2.900000	3.700000	3.266667	

			temp_dailyMedian	hum_dailyMin	hum_dailyMax	
hum_dailyMean	\					
year	month	day				
2009	1	1	-18.777778	71	79	
74.000000		1	-18.777778	71	79	
74.000000		1	-18.777778	71	79	
74.000000						

			wind_dailyMean	wind_dailyMax	wind_dailyMin
precip	\				
year	month	day			
67.000000	2	-18.722222	66	68	
67.000000	2	-18.722222	66	68	
67.000000	...	...	...	...	...
2023	12	30	10.800000	58	78
66.333333	30	10.800000	58	78	
66.333333	31	3.200000	88	96	
91.000000	31	3.200000	88	96	
91.000000	31	3.200000	88	96	
91.000000	31	3.200000	88	96	
2009	1	1	2.000000	2	0.00
	1	1	2.000000	2	0.00
	1	1	2.000000	2	0.00
	2	1	1.666667	3	0.05
	2	1	1.666667	3	0.05
...	...	...	...	...	...
2023	12	30	2.666667	3	0.00
	30	2	2.666667	3	0.00
	31	1	1.000000	2	5.40
	31	1	1.000000	2	5.40
	31	1	1.000000	2	5.40
			precipType		
year	month	day			
2009	1	1	NaN		
	1	1	NaN		
	1	1	NaN		
	2	2	snow		
	2	2	snow		

```

...
2023 12    30      ...
              30      NaN
              31      rain
              31      rain
              31      rain

[16434 rows x 12 columns]

# tests, DO NOT MODIFY OR COPY THIS CELL!!
assert len(agg_daily_data.columns) >= 11, "according to the
instructions, the dataframe should have >= 13 columns"
assert len(agg_daily_data.index.levels) == 3, "according to the
instructions, the dataframe should have a multi-index with 2 levels"
assert len(agg_daily_data) > 4000
assert len(agg_daily_data.columns) >= 11
assert len(agg_daily_data.index.levels) == 3
assert len(agg_daily_data.index.levels[0]) == 15
assert len(agg_daily_data.index.levels[1]) == 12

```

## Merge cyclist and weather datasets

Merge the `agg_daily_data` and `data_cyclists` datasets. Both dataframes should now be on a daily index. Beware that both datasets contain rows that do not appear in the other dataset.

```

def merge_data(weather_df, traffic_df):
    """
    Parameters
    -----
    weather_df: daily weather data frame
    traffic_df: traffic data frame

    Returns
    -----
    merged_data: merged data frame that contains both daily weather
    observations and traffic incidents
    """
    # YOUR CODE HERE
    merged_data = (
        weather_df
        .join(traffic_df[["number", "day_of_week"]], how="outer")
        .sort_index()
    )

    return merged_data

data_merged = merge_data(agg_daily_data, data_cyclists)
data_merged.head(13600)

```

year	month	day	temp_dailyMin	temp_dailyMax	temp_dailyMean	\
2009	1	1	-20.277778	-17.833333	-18.962963	
		1	-20.277778	-17.833333	-18.962963	
		1	-20.277778	-17.833333	-18.962963	
		2	-19.444444	-18.500000	-18.888889	
		2	-19.444444	-18.500000	-18.888889	
...			...	...	...	
2021	5	29	-9.944444	-7.333333	-9.018519	
		30	-11.222222	-10.500000	-10.888889	
		30	-11.222222	-10.500000	-10.888889	
		30	-11.222222	-10.500000	-10.888889	
		31	-10.055556	-7.333333	-8.574074	
hum_dailyMean	\		temp_dailyMedian	hum_dailyMin	hum_dailyMax	
year	month	day				
2009	1	1	-18.777778	71	79	
74.000000		1	-18.777778	71	79	
74.000000		1	-18.777778	71	79	
74.000000		2	-18.722222	66	68	
67.000000		2	-18.722222	66	68	
67.000000			...	...	...	
...			...	...	...	
2021	5	29	-9.777778	35	62	
50.333333		30	-10.944444	63	75	
68.666667		30	-10.944444	63	75	
68.666667		30	-10.944444	63	75	
68.666667		31	-8.333333	47	60	
51.333333						
precip	\		wind_dailyMean	wind_dailyMax	wind_dailyMin	
year	month	day				
2009	1	1	2.000000	2	2	0.00
		1	2.000000	2	2	0.00
		1	2.000000	2	2	0.00

	2	1.666667	3	0	0.05	
	2	1.666667	3	0	0.05	
	...	...	...	...	...	...
2021	5	29	3.000000	3	3	0.10
		30	3.333333	4	3	1.30
		30	3.333333	4	3	1.30
		30	3.333333	4	3	1.30
		31	2.333333	3	2	0.05
year	month	day	precipType	number	day_of_week	
2009	1	1	NaN	NaN	NaN	
		1	NaN	NaN	NaN	
		1	NaN	NaN	NaN	
		2	snow	NaN	NaN	
		2	snow	NaN	NaN	
	...	...	...	...	...	
2021	5	29	rain	2779.0	Saturday	
		30	rain	2122.0	Sunday	
		30	rain	2122.0	Sunday	
		30	rain	2122.0	Sunday	
		31	rain	3930.0	Monday	
[13600 rows x 14 columns]						
# DO NOT MODIFY OR COPY THIS CELL!!						
data_merged = merge_data(agg_daily_data, data_cyclists)						
data_merged						
year	month	day	temp_dailyMin	temp_dailyMax	temp_dailyMean	\
2009	1	1	-20.277778	-17.833333	-18.962963	
		1	-20.277778	-17.833333	-18.962963	
		1	-20.277778	-17.833333	-18.962963	
		2	-19.444444	-18.500000	-18.888889	
		2	-19.444444	-18.500000	-18.888889	
	...	...	...	...	...	
2023	12	30	7.100000	11.900000	9.933333	
		30	7.100000	11.900000	9.933333	
		31	2.900000	3.700000	3.266667	
		31	2.900000	3.700000	3.266667	
		31	2.900000	3.700000	3.266667	

		temp_dailyMedian	hum_dailyMin	hum_dailyMax
hum_dailyMean	\			
year	month	day		

2009	1	1	-18.777778	71	79
74.000000		1	-18.777778	71	79
74.000000		1	-18.777778	71	79
74.000000		2	-18.722222	66	68
67.000000		2	-18.722222	66	68
67.000000		...	...	...	...
...		...	...	...	...
2023	12	30	10.800000	58	78
66.333333		30	10.800000	58	78
66.333333		31	3.200000	88	96
91.000000		31	3.200000	88	96
91.000000		31	3.200000	88	96
91.000000		31	3.200000	88	96
91.000000		...	...	...	...

		wind_dailyMean	wind_dailyMax	wind_dailyMin
precip	\			
year	month	day		

2009	1	1	2.000000	2	0.00
		1	2.000000	2	0.00
		1	2.000000	2	0.00
		2	1.666667	3	0 0.05
		2	1.666667	3	0 0.05
...		...	...	...	...
2023	12	30	2.666667	3	2 0.00
		30	2.666667	3	2 0.00
		31	1.000000	2	0 5.40

31	1.000000	2	0	5.40	
31	1.000000	2	0	5.40	
year	month	day	precipType	number	day_of_week
2009	1	1	NaN	NaN	NaN
		1	NaN	NaN	NaN
		1	NaN	NaN	NaN
		2	snow	NaN	NaN
		2	snow	NaN	NaN
...			...	...	...
2023	12	30	NaN	NaN	NaN
		30	NaN	NaN	NaN
		31	rain	NaN	NaN
		31	rain	NaN	NaN
		31	rain	NaN	NaN

[16434 rows x 14 columns]

## Multiple Choice Questions

Question 1: Create a scatterplot with minimum daily temperature on the x-axis and the number of cyclists on the y-axis. Which of the following statement(s) about the scatterplot is/are correct?

- a) The correlation between minimum daily temperature and the number of cyclists is negative.
- b) The higher the minimum daily temperatur, the more cyclists are counted.
- c) The association between minimum daily temperature and the number of cyclists is u-shaped.
- d) There is no association between minimum daily temperature and the number of cyclists.

```
answer_1 = {"a": 0, "b": 1, "c": 0, "d": 0}

# TEST (will be hidden in lab exam)
assert answer_1["a"] == 0
assert answer_1["b"] == 1
assert answer_1["c"] == 0
assert answer_1["d"] == 0

from scipy.stats import pearsonr
x_col = "temp_dailyMin"    # change if your min temp column has a
                            # different name
y_col = "number"

plot_df = (
    data_merged[[x_col, y_col]]
```


```

    .dropna()
    .rename(columns={x_col: "min_temp_c", y_col: "cyclists"})
)

plt.figure(figsize=(7, 5))
sns.regplot(
    data=plot_df,
    x="min_temp_c",
    y="cyclists",
    scatter_kws={"alpha": 0.4, "s": 25},
    line_kws={"linewidth": 2}
)
plt.xlabel("Minimum daily temperature, °C")
plt.ylabel("Cyclists counted")
plt.title("Cyclists vs minimum daily temperature")
plt.tight_layout()
plt.show()

r, p = pearsonr(plot_df["min_temp_c"], plot_df["cyclists"])
print(f"Pearson r = {r:.3f}, p value = {p:.3g}")

```



Pearson r = 0.665, p value = 3.77e-280

```

# YOUR CODE HERE
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Day of week line plot
order =
["Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday", "Sunday"]

wk = (
    data_merged[["day_of_week", "number"]]
    .dropna()
    .assign(day_of_week=lambda d: pd.Categorical(d["day_of_week"]),
categories=order, ordered=True))
    .groupby("day_of_week", as_index=False)[ "number"].mean()
    .sort_values("day_of_week")
)

plt.figure(figsize=(7, 4.5))
plt.plot(wk[ "day_of_week"], wk[ "number"], marker="o")
plt.xlabel("Day of week")
plt.ylabel("Average cyclists")
plt.title("Average cyclists by day of week")
plt.grid(axis="y", alpha=0.3)
plt.tight_layout()
plt.show()

# Weekly pattern by season, fix uses a single reset frame
tmp = data_merged.reset_index()
[[ "year", "month", "day", "day_of_week", "number"]].copy()


def to_season(m):
    if m in [12, 1, 2]:
        return "winter"
    if m in [3, 4, 5]:
        return "spring"
    if m in [6, 7, 8]:
        return "summer"
    return "autumn"

tmp[ "season"] = tmp[ "month"].map(to_season)
tmp[ "day_of_week"] = pd.Categorical(tmp[ "day_of_week"],
categories=order, ordered=True)


wk_season = (
    tmp.dropna(subset=[ "day_of_week", "number"])
    .groupby([ "season", "day_of_week"], as_index=False)[ "number"]
    .mean()
    .sort_values(["season", "day_of_week"]))

```

```
)  
  
plt.figure(figsize=(8, 5))  
sns.lineplot(data=wk_season, x="day_of_week", y="number",  
hue="season", marker="o")  
plt.xlabel("Day of week")  
plt.ylabel("Average cyclists")  
plt.title("Average cyclists by day of week and season")  
plt.grid(axis="y", alpha=0.3)  
plt.tight_layout()  
plt.show()  
  
/tmp/ipykernel_349/3066950987.py:13: FutureWarning: The default of  
observed=False is deprecated and will be changed to True in a future  
version of pandas. Pass observed=False to retain current behavior or  
observed=True to adopt the future default and silence this warning.  
.groupby("day_of_week", as_index=False)[ "number"].mean()
```



```
/tmp/ipykernel_349/3066950987.py:43: FutureWarning: The default of  
observed=False is deprecated and will be changed to True in a future  
version of pandas. Pass observed=False to retain current behavior or  
observed=True to adopt the future default and silence this warning.  
.groupby(["season", "day_of_week"], as_index=False)[ "number"]
```



Question 2 Make a line plot with the days of the week on the x-axis and the number of cyclists on the y-axis for each month of the year. Which of the following statements is correct?

- a) The number of cyclists is stable at ~2000 throughout the year
- b) During summer months (June–August) the number of cyclists on weekends is higher than on workdays.
- c) In January, the number of cyclists is much higher at the beginning of the workweek (Monday) than at the end (Friday)
- d) The number of cyclists is lower during the winter months, and higher during the summer.

```
answer_2 = ["a": 0, "b": 0, "c": 0, "d": 0]

assert answer_2["a"] == 0
assert answer_2["b"] == 0
assert answer_2["c"] == 0
assert answer_2["d"] == 0

# YOUR CODE HERE

import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

# Choose candidates, adjust names if needed
```

```

candidates = {
    "day_type": "day_type",           # expected values like weekday,
weekend
    "temp_dailyMean": "temp_dailyMean", # numeric month one to twelve
    "month": "month",                 # daily precipitation
    "precip": "precip",
}
df = data_merged.reset_index().copy()

# If your frame has day_of_week but not day_type, derive a simple
day_type
if "day_type" not in df.columns and "day_of_week" in df.columns:
    df["day_type"] =
df["day_of_week"].isin(["Saturday", "Sunday"]).map({True:"weekend",
False:"weekday"})

use_cols = [c for c in candidates.values() if c in df.columns]
model_df = df[use_cols + ["number"]].dropna()

# One hot encode categoricals, keep numeric as is
X = pd.get_dummies(
    model_df.drop(columns=["number"]),
    columns=[c for c in ["day_type", "month"] if c in
model_df.columns],
    drop_first=True
)
y = model_df["number"].values

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25, random_state=42)

lr = LinearRegression()
lr.fit(X_train, y_train)
pred = lr.predict(X_test)

print("R2 on test set:", round(r2_score(y_test, pred), 3))
coefs = pd.Series(lr.coef_, index=X.columns).sort_values(key=abs,
ascending=False)
print("Top coefficients:")
print(coefs.head(10))

R2 on test set: 0.761
Top coefficients:
day_type_weekend      -1290.776473
month_9                  1080.874603
month_10                  990.044502
month_6                   925.749166
month_11                  849.577386
month_5                   776.361427

```

```
month_3           693.773096
month_7           485.615145
month_8           393.459603
month_4           334.730035
dtype: float64
```

Question 3 Which of the variables in the dataset are a suitable dependent variable in a classification task?

- a) day\_of\_week, when dichotomised into two categories (workweek and weekend)
- b) temp\_dailyMean
- c) month
- d) precip

```
answer_3 = {"a": 0, "b": 0, "c": 0, "d": 0}

assert answer_3["a"] == 1
assert answer_3["b"] == 0
assert answer_3["c"] == 1
assert answer_3["d"] == 0
```

Question 4 When the model performs well on the training data and poorly on new, unseen data, what is this indicative of?

- a) Overfitting
- b) Underfitting
- c) High bias
- d) High variance

```
answer_4 = {"a": 0, "b": 0, "c": 0, "d": 0}

assert answer_4["a"] == 1txt
assert answer_4["b"] == 0
assert answer_4["c"] == 0
assert answer_4["d"] == 1
```